

# Sockets

## **Advanced Embedded Software Development**

with **Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

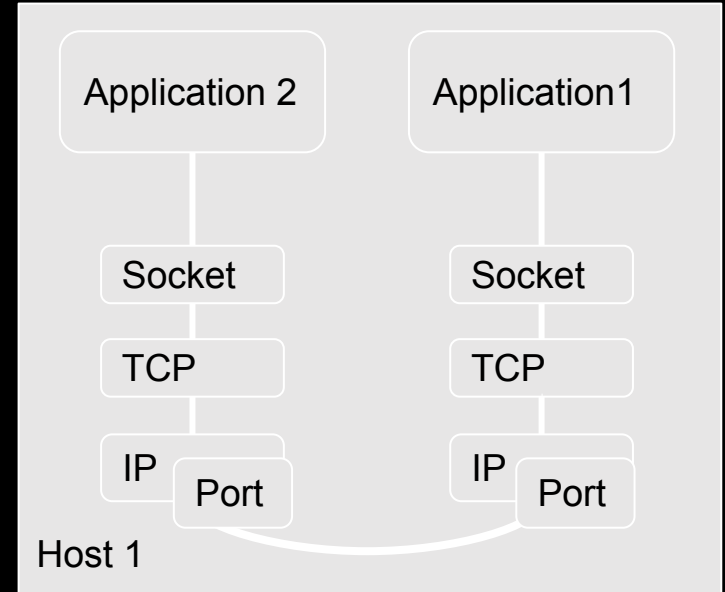

Understand Linux Sockets

Understand How to Use Sockets in your programs

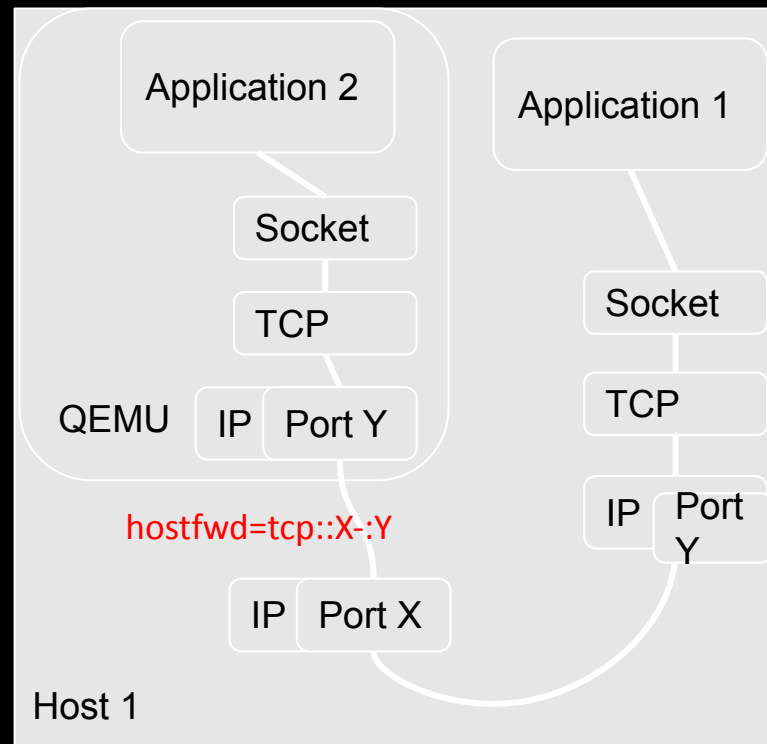
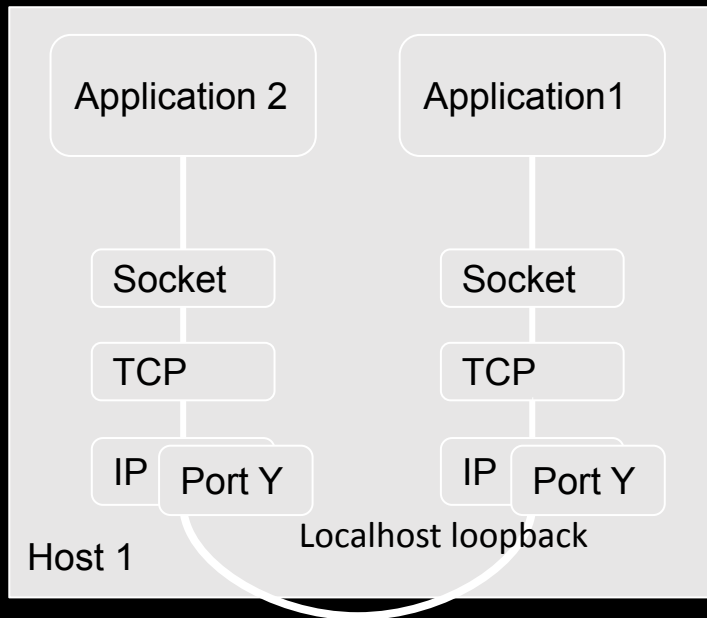
# Sockets

- One of several forms of Interprocess Communication (IPC)
  - Can communicate across different systems over TCP/IP
- Also known as BSD or Berkeley Sockets
- Supported by all major operating systems

# Connection Oriented Sockets



# Port Forwarding



- Allows us to route port to VM or emulator

# TCP

- Transmission Control Protocol
- Connection oriented protocol
  - Connection is established and maintained while programs are exchanging messages
- Accepts packets
- Manages flow control
- Handles retransmission of dropped packets
- Handles acknowledgement of packets

# IP

- Addresses Packets
- Supports routing between sender and receiver
- IPv4 was first implementation - X.X.X.X - 4 byte (32 bit) format - ~4.3 billion addresses
- IPv6 supports 128 bits of address space - 340 billion billion billion billion addresses)
  - 4000 addresses for each person on earth
- Also uses a “port” for local addressing

# Types of Sockets

- SOCK\_STREAM - Stream Sockets
  - Reliable two way connected (TCP) streams
  - Messages are delivered in order
  - Retried as necessary
- SOCK\_DGRAM - Datagram Sockets
  - Connectionless sockets
  - Use UDP (User Datagram Protocol) instead of TCP



# Accessing Sockets

- In Linux everything is a file
- How do we interact with sockets?
  - Using a socket file descriptor
- How do we obtain a socket file descriptor?
  - Use the `socket()` POSIX function

## NAME

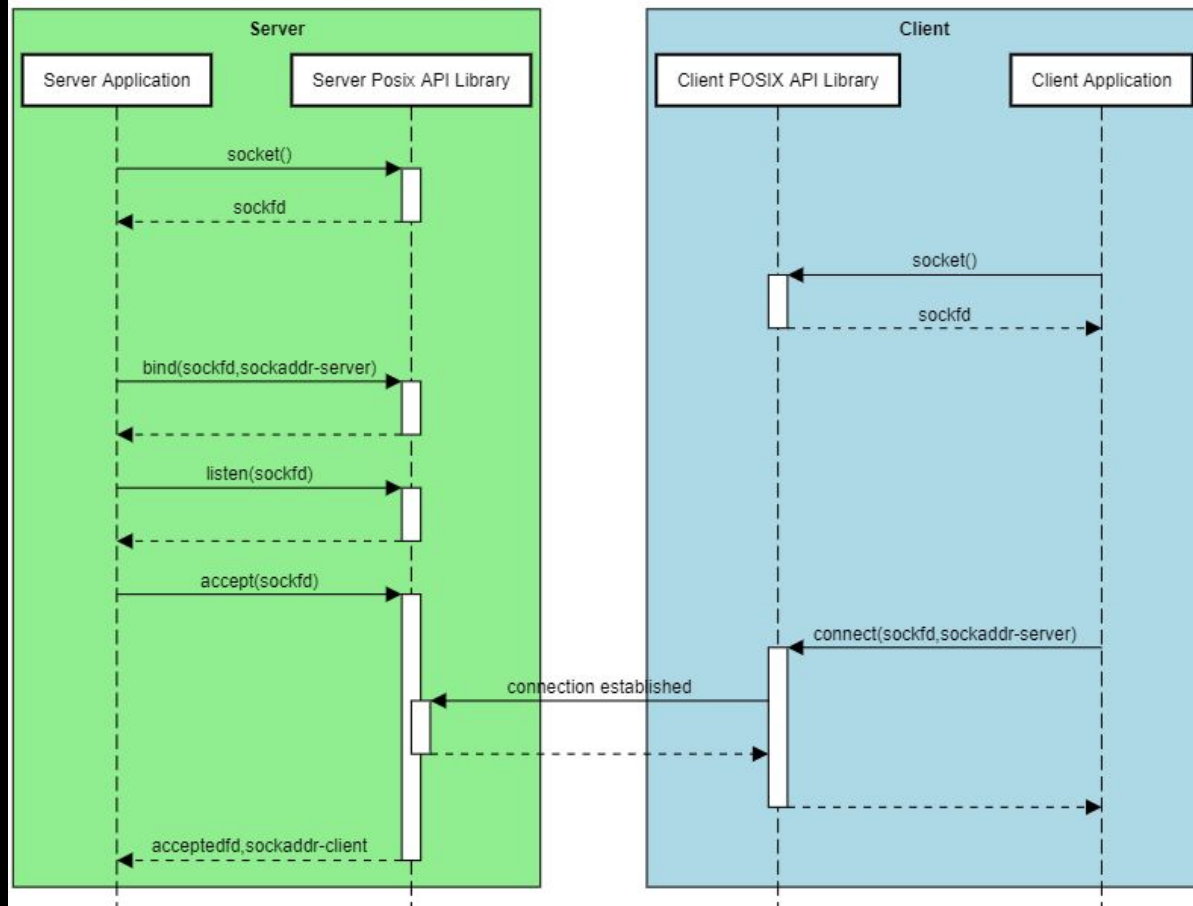
`socket` - create an endpoint for communication

## SYNOPSIS

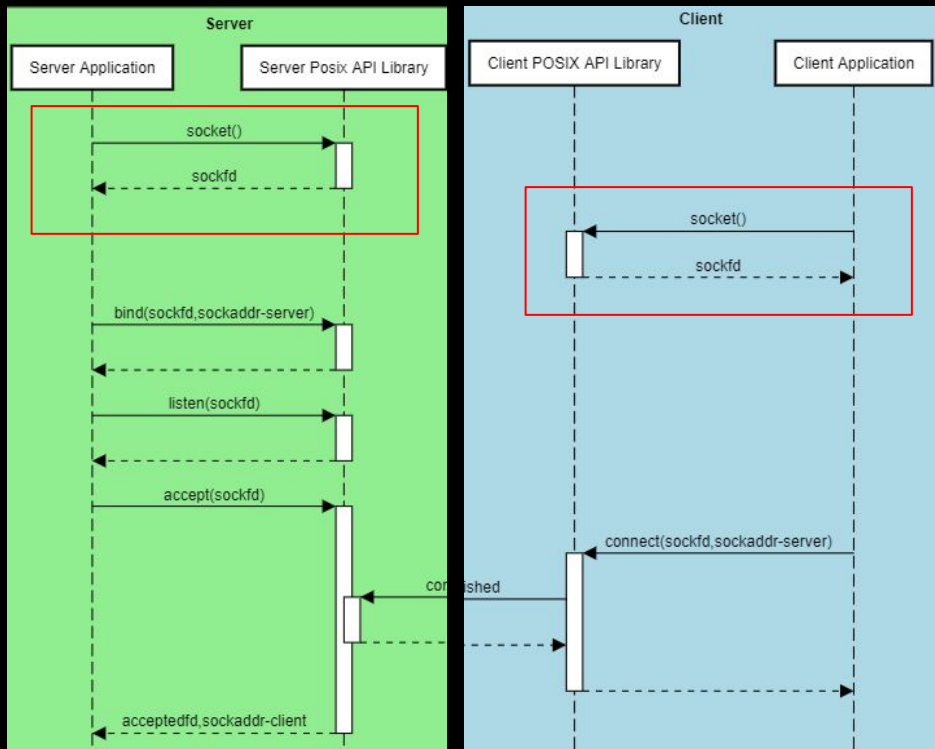
```
#include <sys/types.h>           /* See NOTES */
#include <sys/socket.h>

int socket(int domain, int type, int protocol);
```

## Socket Creation



# Socket



## NAME

`socket` - create an endpoint for communication

## SYNOPSIS

```
#include <sys/types.h>           /* See NOTES */
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

- `domain` - `PF_INET` or `PF_INET6`
- `type` `SOCK_STREAM` or `SOCK_DGRAM`
- `protocol` - 0 to choose proper for given type

# bind

NAME  
bind - bind a name to a socket

SYNOPSIS  
`#include <sys/types.h> /* See NOTES */`  
`#include <sys/socket.h>`  
  
`int bind(int sockfd, const struct sockaddr *addr,  
socklen_t addrlen);`

- Assigns an address to the socket
- sockfd is the fd from socket()
- sockaddr addr/addrlen describes the address to bind (optionally select a specific network adapter).

# bind - sockaddr

```
NAME
    bind - bind a name to a socket

SYNOPSIS
#include <sys/types.h>          /* See NOTES */
#include <sys/socket.h>

int bind(int sockfd, const struct sockaddr *addr,
         socklen_t addrlen);

struct sockaddr {
    sa_family_t sa_family;
    char        sa_data[14];
}
```

- sockaddr - maps to a server socket location
  - sa\_family - AF\_INET or AF\_INET6
  - sa\_data - destination address and port
- socketlen\_t - unsigned integer type  
sizeof(struct sockaddr)

# Setting up sockaddr

- sockaddr structure isn't built directly, setup from other structures
- Two options discussed:
  - Setting up sockaddr\_in and casting to sockaddr
  - Setting up with getaddrinfo()
  - getaddrinfo is newer and more flexible
  - Either is acceptable for the assignment

# Setting up sockaddr

```
struct sockaddr {
    unsigned short    sa_family;    // address family, AF_XXX
    char              sa_data[14];  // 14 bytes of protocol address
};
```

```
struct addrinfo {
    int                ai_flags;     // AI_PASSIVE, AI_CANONNAME, etc.
    int                ai_family;    // AF_INET, AF_INET6, AF_UNSPEC
    int                ai_socktype;  // SOCK_STREAM, SOCK_DGRAM
    int                ai_protocol;  // use 0 for "any"
    size_t             ai_addrlen;   // size of ai_addr in bytes
    struct sockaddr *ai_addr;        // struct sockaddr_in or _in6
    char               ai_canonname; // full canonical hostname

    struct addrinfo *ai_next;        // linked list, next node
};
```

## NAME

getaddrinfo, freeaddrinfo, gai\_strerror - network address and service translation

## SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
```

```
int getaddrinfo(const char *node, const char *service,
                const struct addrinfo *hints,
                struct addrinfo **res);
```

```
void freeaddrinfo(struct addrinfo *res);
```

- getaddrinfo provides addrinfo (containing sockaddr) through addrinfo arguments

# Setting addrinfo hints and node

```
struct addrinfo {  
    int      ai_flags;    // AI_PASSIVE, AI_CANONNAME, etc.  
    int      ai_family;   // AF_INET, AF_INET6, AF_UNSPEC  
    int      ai_socktype; // SOCK_STREAM, SOCK_DGRAM  
    int      ai_protocol; // use 0 for "any"  
    size_t   ai_addrlen;  // size of ai_addr in bytes  
    struct sockaddr *ai_addr; // struct sockaddr_in or _in6  
    char      *ai_canonname; // full canonical hostname  
  
    struct addrinfo *ai_next; // linked list, next node  
};
```

```
int getaddrinfo(const char *node, const char *service,  
               const struct addrinfo *hints,  
               struct addrinfo **res);
```

If the `AI_PASSIVE` flag is specified in `hints.ai_flags`, and `node` is `NULL`, then the returned socket addresses will be suitable for `bind(2)`ing a socket that will `accept(2)` connections.

- `ai_flags` in `hints` and `node` parameter sets up the socket address for `bind()/accept()`
  - `hints.ai_flags = AI_PASSIVE`, `hints.ai_family = AF_INET`, `hints.ai_socktype = SOCK_STREAM`
  - `node = NULL`



# Setting getaddrinfo service

```
int getaddrinfo(const char *node, const char *service,  
               const struct addrinfo *hints,  
               struct addrinfo **res);
```

service sets the port in each returned address structure. If this argument is a service name (see `services(5)`), it is translated to the corresponding port number. This argument can also be specified as a decimal number, which is simply converted to binary. If service is NULL,

- service parameter sets port for the connection
  - “1234” would setup for port 1234

# Getting sockaddr

```
int getaddrinfo(const char *node, const char *service,  
               const struct addrinfo *hints,  
               struct addrinfo **res);
```

- Setup pointer res to store addrinfo returned from getaddrinfo
  - Pass address of pointer as res arg (pointer to pointer)

# Getting sockaddr - pointer to pointer

```
int status;
struct addrinfo hints;
struct addrinfo *servinfo; // will point to the results

memset(&hints, 0, sizeof hints); // make sure the struct is empty
hints.ai_family = AF_UNSPEC;      // don't care IPv4 or IPv6
hints.ai_socktype = SOCK_STREAM; // TCP stream sockets
hints.ai_flags = AI_PASSIVE;     // fill in my IP for me

if ((status = getaddrinfo(NULL, "3490", &hints, &servinfo)) != 0) {
    fprintf(stderr, "getaddrinfo error: %s\n", gai_strerror(status));
    exit(1);
}

// servinfo now points to a linked list of 1 or more struct addrinfos

// ... do everything until you don't need servinfo anymore ....

freeaddrinfo(servinfo); // free the linked-list
```

```
int getaddrinfo(const char *node,      // e.g. "www.example.com" or IP
               const char *service,  // e.g. "http" or port number
               const struct addrinfo *hints,
               struct addrinfo **res);
```

```
struct addrinfo {
    int      ai_flags;      // AI_PASSIVE, AI_CANONNAME, etc.
    int      ai_family;    // AF_INET, AF_INET6, AF_UNSPEC
    int      ai_socktype;  // SOCK_STREAM, SOCK_DGRAM
    int      ai_protocol;  // use 0 for "any"
    size_t   ai_addrlen;   // size of ai_addr in bytes
    struct sockaddr *ai_addr; // struct sockaddr_in or _in6
    char      *ai_canonname; // full canonical hostname

    struct addrinfo *ai_next; // linked list, next node
};
```

passed by  
reference to  
getaddrinfo

allocated on  
stack

malloc'd  
inside  
getaddrinfo,  
assigned to  
&servinfo

&servinfo

struct  
addrinfo \*  
servinfo


struct  
addrinfo

# Getting sockaddr

```
int getaddrinfo(const char *node, const char *service,  
               const struct addrinfo *hints,  
               struct addrinfo **res);
```

- Call getaddrinfo with hints, port string as service argument, and pointer to pointer in res
- Use res->ai\_addr as sockaddr for bind()
- freeaddrinfo(res) when no longer needed
  - What if you forget to free?
    - Memory Leak

# listen



## NAME

`listen` - listen for connections on a socket

## SYNOPSIS

```
#include <sys/types.h>           /* See NOTES */
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

- Passed sockfd from socket()
- backlog specifies number of pending connections allowed before refusing

# accept



NAME

accept, accept4 - accept a connection on a socket

SYNOPSIS


```
#include <sys/types.h>
#include <sys/socket.h>
```

/\* See NOTES \*/

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- sockfd - socket file descriptor from socket()
- addr - location to store the connecting address
- addrlen in/out - length of addr and location to store the length of result
- returns: fd for accepted connection

<https://stackoverflow.com/questions/27014955/socket-connect-vs-bind>  
<http://man7.org/linux/man-pages/man2/accept.2.html>



# recv/send

- Similar to read/write file descriptor based commands we've discussed in early lectures

## RECV(2)

### NAME

recv, recvfrom, recvmsg - receive a message from a socket

### SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

## SEND(2)

### NAME

send, sendto, sendmsg - send a message on a socket

### SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

## READ(2)

### NAME

read - read from a file descriptor

### SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

## WRITE(2)

### NAME

write - write to a file descriptor

### SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

<https://stackoverflow.com/questions/27014955/socket-connect-vs-bind>

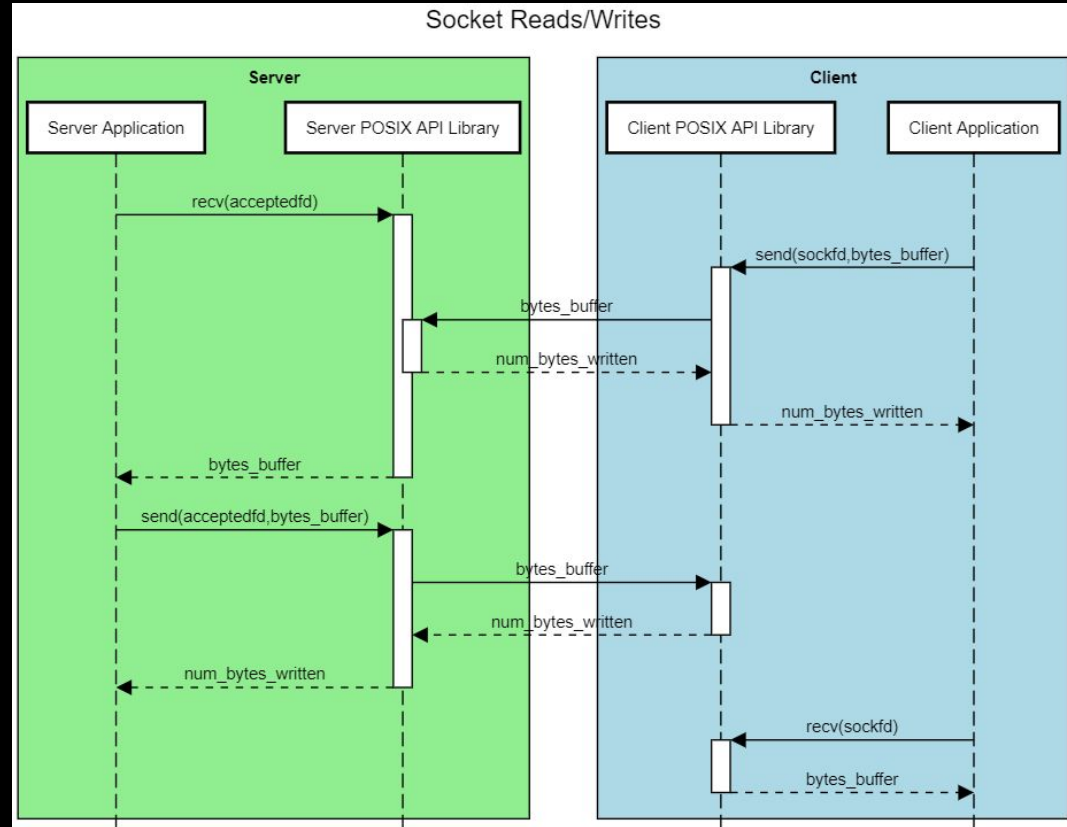
<http://man7.org/linux/man-pages/man2/read.2.html>

<http://man7.org/linux/man-pages/man2/write.2.html>

# recv/send



- Use `acceptedfd` (from `accept()` return value) for server application
- Use blocking or non-blocking reads based on flags argument to `recv/send`



<https://stackoverflow.com/questions/27014955/socket-connect-vs-bind>

<http://man7.org/linux/man-pages/man2/read.2.html>

<http://man7.org/linux/man-pages/man2/write.2.html>