

Introduction to Kernel Modules

**Advanced Embedded Linux
Development**
with **Dan Walkes**



University of Colorado **Boulder**

Learning objectives:

Kernel Module Layout and Content
Comparison with Application
Development

Hello World Driver

- Regarding “Hello World” in Chapter 2:
 - Use repository <https://github.com/cu-ecen-5013/ldd3-samples> which has been updated to support latest kernels.
 - You can use the kernel in your VM to build, no need to download from kernel.org
 - Run make from the misc-modules directory
 - insmod as root
 - The output of hello world will be in /var/log/syslog

```
Aug 17 15:40:03 ecen5013-VirtualBox kernel: [105871.303956] hello: loading out-of-tree module taints kernel.  
Aug 17 15:40:03 ecen5013-VirtualBox kernel: [105871.305224] Hello, world  
Aug 17 15:40:49 ecen5013-VirtualBox kernel: [105916.782604] Goodbye, cruel world
```

Tainted Kernel

```
Aug 17 15:40:03 ecen5013-VirtualBox kernel: [105871.303956] hello: loading out-of-tree module taints kernel.  
Aug 17 15:40:03 ecen5013-VirtualBox kernel: [105871.305224] Hello, world  
Aug 17 15:40:49 ecen5013-VirtualBox kernel: [105916.782604] Goodbye, cruel world
```

- What does this mean?
 - Reference regarding non GPL module tainting in Chapter 2.
 - Other reasons for taint including “out-of-tree module”
 - Signal to kernel developers that they don’t have the information to debug a bug report.

hello.c

```
/*
 * $Id: hello.c,v 1.5 2004/10/26 03:32:21 corbet Exp $
 */
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");

static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world\n");
    return 0;
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Module vs Application

- `module_init(funcname)` identifies a function called when the module starts (from `insmod`)
 - Same as `int main(void)`?
 - Different in that it sets up module functions and exits.
 - Similar to an event driven application.

Module vs Application

- `module_exit(funcname)` identifies a function called when the module is unloaded (from `rmmod`)
- Similar to signal handler for SIGTERM.
 - Difference with application programming:
No automatic cleanup when process exits
 - If you forget to free memory it stays allocated after your driver is unloaded.

Module vs Application

- Why printk and not printf?
 - We don't have libc!
 - No <stdio.h>!

Module vs Application

- `MODULE_LICENSE`
 - Clarifies that the module bears a free license.
 - Some kernel function calls aren't available with proprietary licenses.
- Segfault kills more than your code most likely, may bring down the entire system.

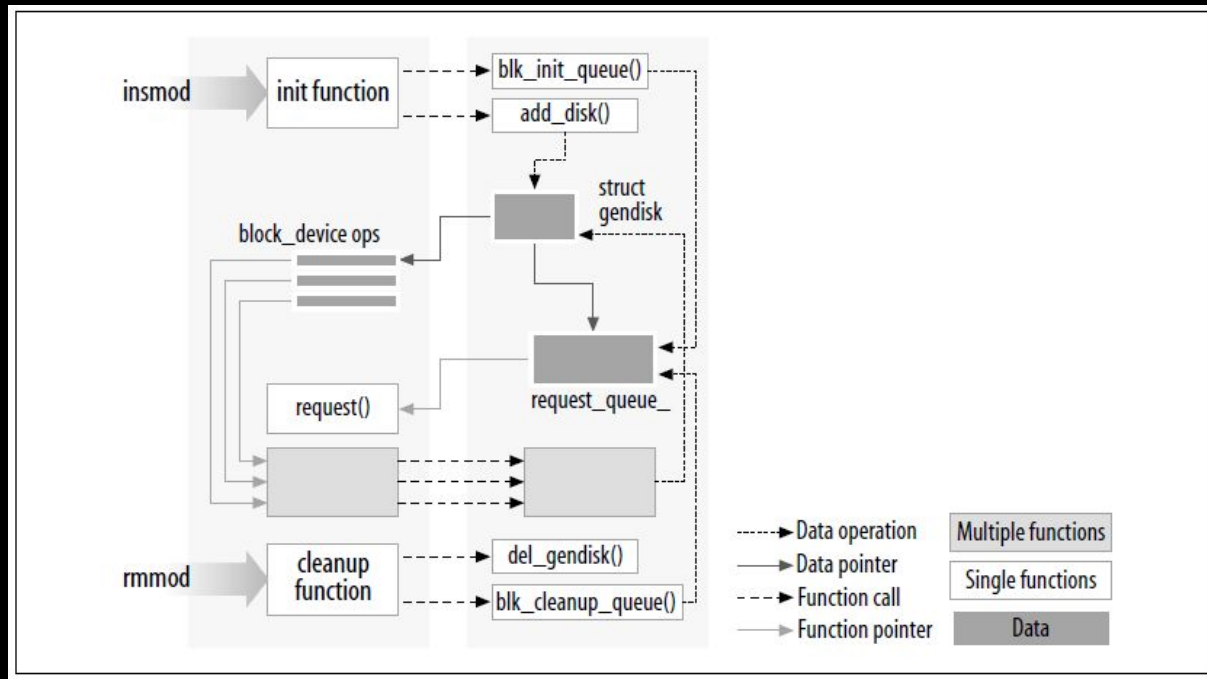
Module vs Application

- Kernel stack is small compared to applications
 - Application default stack size is ~2 MB in size
 - Kernel stack may be 4k in size

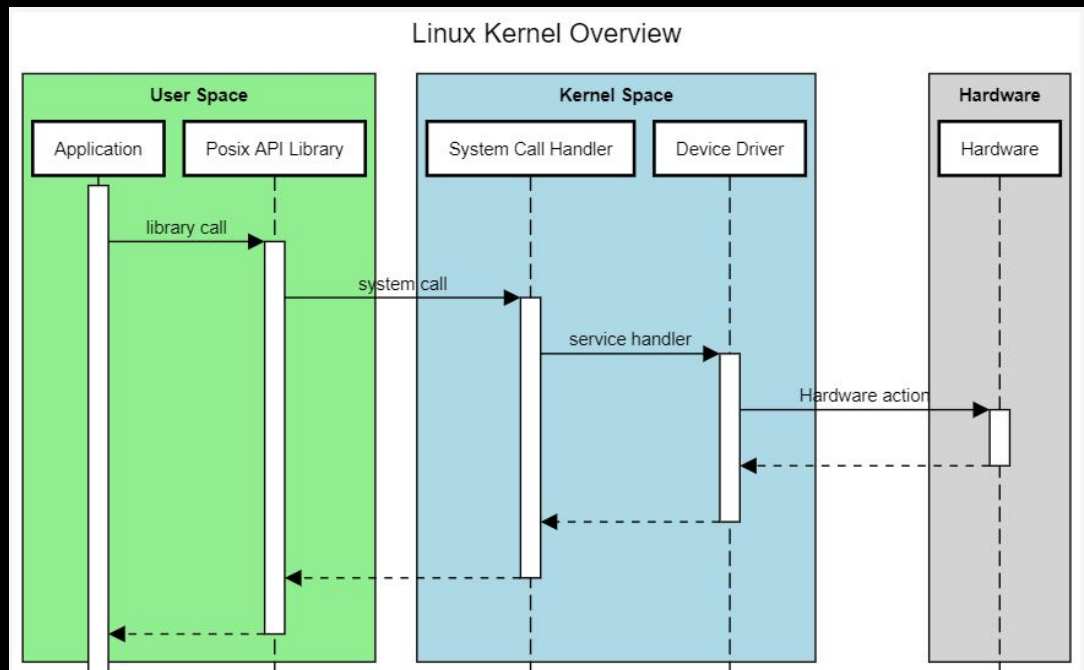
Module vs Application

- Kernel stack may be 4k in size
 - All active functions (including yours) share this small space.
 - Don't declare large automatic stack allocated variables
 - Instead allocate memory dynamically
- Floating point is generally not supported

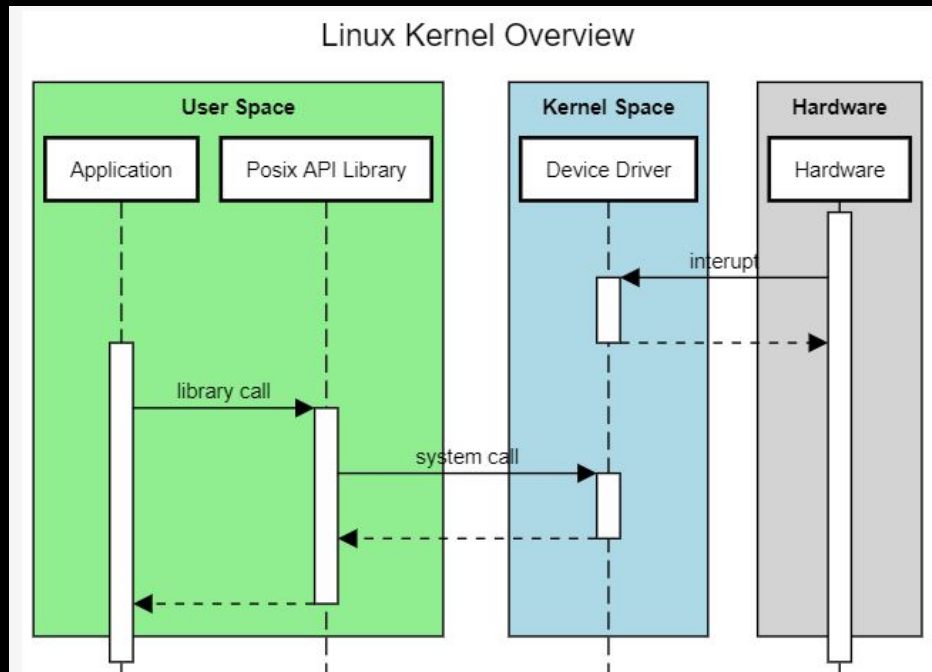
Linking a module into the kernel



Kernel entry - syscalls



Kernel Entry - Interrupts



User Space vs Kernel Space - Resource Access

- Modern processors support enforcing protection against unauthorized access to resources.
 - Operating modalities or levels
 - Separate address space

User Space vs Kernel Space - Resource Access

- Applications Use User Space
 - Executes at lowest privileged operating level
 - Access to hardware is regulated
 - Access to memory is regulated

User Space vs Kernel Space - Resource Access

- Drivers Use Kernel Space
 - Executes at highest privileged operating level (supervisor mode)
 - Entered with system call or hardware interrupt
 - Your module may handle one or both

User Space vs Kernel Space - Concurrency

- The Kernel runs multiple processes, multiple processes may be trying to use your driver.
- Interrupts run asynchronously (and kernel timers)
- SMP (symmetric multiprocessor) systems are supported
 - Two or more processors running on the same memory and OS instance

User Space vs Kernel Space - Concurrency

- Kernel code (including your driver code) must be reentrant
 - Data structures must keep threads of execution separate
 - Shared data access must be handled correctly
- Reentrancy must be considered, even when not sleeping