

Connecting to Userspace

**Advanced Embedded Linux
Development**
with **Dan Walkes**



University of Colorado **Boulder**

Learning objectives:

Creating and using device nodes

Registering and Unregistering Devices

Linux Device Driver Examples

- “scull” driver overview
 - “Simple Character Utility for Loading Localities”
 - Acts on memory as if it were a device
- Benefits
 - Portable across architectures
 - Not hardware dependent
- Drawbacks
 - Not a “real” device driver
- See <https://github.com/cu-ecen-5013/ldd3>

Types of “scull” drivers

- scull0 to scull3
 - Four devices with memory array global and persistent
 - Global: data shared with all file descriptors
 - Persistent: Not lost if device is closed and reopened.
- scullpipe0 to scullpipe3
 - FIFO pipes
 - Implement blocking and non blocking read and write.

Types of “scull” drivers

- `scullsingle`
 - `scull` but only allows use for a single process
- `scullpriv`
 - `scull` private to each virtual console session
- `sculluid` and `scullwuid`
 - Opened only by one user at a time.

/dev Directory Setup

- How do devices in the /dev directory map to appropriate kernel modules?
 - Device numbers/mknod
 - From Assignment 3:
- `mknod <name> <type> <major> <minor>`
 - Null device is a known major 1 minor 3
 - Console device is known major 5 minor 1

```
sudo mknod -m 666 dev/null c 1 3
```

Device Numbers

- `ls -l` on the `/dev` directory shows two numbers
 - These are the “major” and “minor” device number
 - major is type - identifies the driver
 - minor identifies a specific device

```

crw-rw-rw-    1 root    root      1,   3 Apr 11  2002 null
crw-----    1 root    root     10,   1 Apr 11  2002 psaux
crw-----    1 root    root      4,   1 Oct 28 03:04 tty1
crw-rw-rw-    1 root    tty       4,  64 Apr 11  2002 ttys0
crw-rw----    1 root    uucp      4,  65 Apr 11  2002 ttyS1
crw--w----    1 vcsa    tty       7,   1 Apr 11  2002 vcs1
crw--w----    1 vcsa    tty      7, 129 Apr 11  2002 vcsa1
crw-rw-rw-    1 root    root      1,   5 Apr 11  2002 zero

```

Common major/minor numbers

- See <https://github.com/torvalds/linux/blob/master/Documentation/admin-guide/devices.txt>
 - Common device drivers use have dedicated major numbers (and sometimes minor numbers)
 - Everyone else allocates

```
8 block      SCSI disk devices (0-15)
  0 = /dev/sda      First SCSI disk whole disk
 16 = /dev/sdb      Second SCSI disk whole disk
 32 = /dev/sdc      Third SCSI disk whole disk
  ...
240 = /dev/sdp      Sixteenth SCSI disk whole disk
```

```
ecen5013@ecen5013-VirtualBox:~$ ls -la /dev/sd*
brw-rw---- 1 root disk 8,  0 Sep 2 08:21 /dev/sda
brw-rw---- 1 root disk 8,  1 Sep 2 08:21 /dev/sda1
brw-rw---- 1 root disk 8,  2 Sep 2 08:21 /dev/sda2
brw-rw---- 1 root disk 8,  5 Sep 2 08:21 /dev/sda5
brw-rw---- 1 root disk 8, 16 Sep 2 08:21 /dev/sdb
brw-rw---- 1 root disk 8, 17 Sep 2 08:21 /dev/sdb1
```


Register/Unregister Example

- See the “sleepy” module for a simple example

```
int sleepy_init(void)
{
    int result;

    /*
     * Register your major, and accept a dynamic number
     */
    result = register_chrdev(sleepy_major, "sleepy", &sleepy_fops);
    if (result < 0)
        return result;
    if (sleepy_major == 0)
        sleepy_major = result; /* dynamic */
    return 0;
}

void sleepy_cleanup(void)
{
    unregister_chrdev(sleepy_major, "sleepy");
}

module_init(sleepy_init);
module_exit(sleepy_cleanup);
```

unregister_chrdev

- What happens to a chrdev I register when my module exits?
 - You are responsible for unregistering
 - Use `unregister_chrdev` for this

```
/**
 * __unregister_chrdev - unregister and destroy a cdev
 * @major: major device number
 * @baseminor: first of the range of minor numbers
 * @count: the number of minor numbers this cdev is occupying
 * @name: name of this range of devices
 *
 * Unregister and destroy the cdev occupying the region described by
 * @major, @baseminor and @count. This function undoes what
 * __register_chrdev() did.
 */
void __unregister_chrdev(unsigned int major, unsigned int baseminor,
                        unsigned int count, const char *name)
```

```
static inline void unregister_chrdev(unsigned int major, const char *name)
{
    __unregister_chrdev(major, 0, 256, name);
}
```

Device Number Representation

- dev_t type holds both, individual major or minor values accessed with macros
- Just a typedef for a 32 bit integer

```
#ifndef u32
#define u32 unsigned int
#endif
```

```
typedef u32 __kernel_dev_t;

typedef __kernel fd_set      fd_set;
typedef __kernel dev_t      dev_t;
```

Device Number Representation

- Create a dev_t:

```
MKDEV(int major, int minor);
```

```
#define MINORBITS      20  
#define MINORMASK      ((1U << MINORBITS) - 1)  
  
#define MAJOR(dev)      ((unsigned int) ((dev) >> MINORBITS))  
#define MINOR(dev)      ((unsigned int) ((dev) & MINORMASK))  
#define MKDEV(ma,mi)    (((ma) << MINORBITS) | (mi))
```

- Extract major and minor numbers from dev_t:

```
MAJOR(dev_t dev);  
MINOR(dev_t dev);
```

Register/Unregister Example

- See the “scull” module for a more complex example
 - Register a range of devices
 - Handle errors after registration

```
int scull_init_module(void)
{
    int result, i;
    dev_t dev = 0;

    /*
     * Get a range of minor numbers to work with, asking for a dynamic
     * major unless directed otherwise at load time.
     */
    if (scull_major) {
        dev = MKDEV(scull_major, scull_minor);
        result = register_chrdev_region(dev, scull_nr_devs, "scull");
    } else {
        result = alloc_chrdev_region(&dev, scull_minor, scull_nr_devs,
                                     "scull");
        scull_major = MAJOR(dev);
    }
    if (result < 0) {
        printk(KERN_WARNING "scull: can't get major %d\n", scull_major);
        return result;
    }

    /*
     * allocate the devices -- we can't have them static, as the number
     * can be specified at load time
     */
    scull_devices = kmalloc(scull_nr_devs * sizeof(struct scull_dev), GFP_KERNEL);
    if (!scull_devices) {
        result = -ENOMEM;
        goto fail; /* Make this more graceful */
    }

    return 0; /* succeed */

fail:
    scull_cleanup_module();
    return result;
}
```

/proc/devices

- Contains current list of allocated devices and associated driver
 - Dynamically registered devices will show up here with name specified in register_chrdev and allocated major number

```
ecen5013@ecen5013-VirtualBox:~$ cat /proc/devices
Character devices:
 1 mem

Block devices:
 7 loop
 8 sd
 9 md
```

Dynamically assigning Devices

- Registering devices will allocate major/minor combinations we can use with `mknod` to create `/dev` file system entries.
- Can't create device nodes (with `mknod`) in advance when using dynamic allocation.
- **How can we handle dynamic allocation?**
 - Parse `/proc/devices` after loading the module to find major number.
 - See `misc-modules module_load` for an example.
 - Other things to do at init time:
 - Set owner user and group
 - Set permissions on device

AWK

- Not an acronym - Abbreviation of developers names
- Works great when you have text output separated by a common field (for instance space)

```
major=$(awk '\$2=="$module" {print \$1}' /proc/devices)
```

```
Block devices:
7 loop
8 sd
```

- `$2 == "sd"` matches when the second field is sd
- `print $1` prints the first field in this case.

Dynamic Device Load Script Example

```
echo "Load our module, exit on failure"
insmod ./module.ko $* || exit 1
echo "Get the major number (allocated with allocate_chrdev_region) from /proc/devices"
major=$(awk "\$2==\"$module\" {print \$1}" /proc/devices)
if [ ! -z ${major} ]; then
    echo "Remove any existing /dev node for /dev/${device}"
    rm -f /dev/${device}
    echo "Add a node for our device at /dev/${device} using mknod"
    mknod /dev/${device} c $major 0
    echo "Change group owner to ${group}"
    chgrp $group /dev/${device}
    echo "Change access mode to ${mode}"
    chmod $mode /dev/${device}
else
    echo "No device found in /proc/devices for driver ${module} (this driver may not allocate a device)"
fi
```

```
aesd@aesd-VirtualBox:~/l3dd3/misc-modules$ sudo ./module_load sleepy
Load our module, exit on failure
Get the major number (allocated with allocate_chrdev_region) from /proc/devices
Remove any existing /dev node for /dev/sleepy
Add a node for our device at /dev/sleepy using mknod
Change group owner to staff
Change access mode to 664
```

```
aesd@aesd-VirtualBox:~/l3dd3/misc-modules$ cat /proc/devices | grep sleepy
241 sleepy
```

```
aesd@aesd-VirtualBox:~/l3dd3/misc-modules$ ls -la /dev | grep sleepy
crw-rw-r-- 1 root staff 241, 0 Feb 29 20:20 sleepy
```

Example Output

Allocated major 241

Device node created
for 241 major, 0 minor