

Additional Kernel Locking Options

**Advanced Embedded Linux
Development**
with **Dan Walkes**



University of Colorado **Boulder**

Learning objectives:

Introduce Kernel Mutexes,
Reader/Writer Semaphores, and
Completions

Mutex()

- Since the book was written, the kernel has added a mutex type for the purposes used in scull

- See

<https://elixir.bootlin.com/linux/v5.3.6/source/Documentation/locking/mutex-design.rst> and <http://lwn.net/Articles/164802/>

Acquire the mutex, uninterruptible:

```
void mutex_lock(struct mutex *lock);
void mutex_lock_nested(struct mutex *lock, unsigned int subclass);
int mutex_trylock(struct mutex *lock);
```

Unlock the mutex:

```
void mutex_unlock(struct mutex *lock);
```

Test if the mutex is taken:

```
int mutex_is_locked(struct mutex *lock);
```

Acquire the mutex, interruptible:

```
int mutex_lock_interruptible_nested(struct mutex *lock,
                                   unsigned int subclass);
int mutex_lock_interruptible(struct mutex *lock);
```

Mutex() map to Semaphore

Acquire the mutex, uninterruptible:

```
void mutex_lock(struct mutex *lock);  
void mutex_lock_nested(struct mutex *lock, unsigned int subclass);  
int mutex_trylock(struct mutex *lock);
```

Unlock the mutex:

```
void mutex_unlock(struct mutex *lock);
```

Test if the mutex is taken:

```
int mutex_is_locked(struct mutex *lock);
```

Acquire the mutex, interruptible:

```
int mutex_lock_interruptible_nested(struct mutex *lock,  
                                   unsigned int subclass);  
int mutex_lock_interruptible(struct mutex *lock);
```

- mutex_lock -> down
- mutex_trylock -> down_trylock
- mutex_lock_interruptible -> down_interruptible
- mutex_unlock -> up
- nested -> used for multiple locks, ordering between

Mutex use in scull

```
struct scull_dev {  
    struct scull_qset *data; /* Pointer to first quantum set */  
    int quantum;             /* the current quantum size */  
    int qset;                /* the current array size */  
    unsigned long size;      /* amount of data stored here */  
    unsigned int access_key; /* used by sculluid and scullpriv */  
    struct mutex lock;       /* mutual exclusion semaphore */  
    struct cdev cdev;        /* Char device structure */  
};
```

- Which initialization method (DEFINE_MUTEX or mutex_init()) would you guess scull uses?
 - mutex_init() runtime

Statically define the mutex:
DEFINE_MUTEX(name);

Dynamically initialize the mutex:
mutex_init(&mutex);

Mutex Use in Scull

```
/* Initialize each device. */
for (i = 0; i < scull_nr_devs; i++) {
    scull_devices[i].quantum = scull_quantum;
    scull_devices[i].qset = scull_qset;
    mutex_init(&scull_devices[i].lock);
    scull_setup_cdev(&scull_devices[i], i);
}
```

```
/*
 * Set up the char_dev structure for this device.
 */
static void scull_setup_cdev(struct scull_dev *dev, int index)
{
    int err, devno = MKDEV(scull_major, scull_minor + index);

    cdev_init(&dev->cdev, &scull_fops);
    dev->cdev.owner = THIS_MODULE;
    dev->cdev.ops = &scull_fops;
    err = cdev_add(&dev->cdev, devno, 1);
    /* Fail gracefully if need be */
    if (err)
        printk(KERN_NOTICE "Error %d adding scull%d", err, index);
}
```

- What does `mutex_init` necessarily happen before?
 - Notifying the kernel about the object in `cdev_add()`
 - “No object can be made available to the kernel before it can function properly”
 - Happens during `module_init()` function

Handling Mutex Lock Failures

```
ssize_t scull_read(struct file *filp, char __user *buf, size_t count,  
                  loff_t *f_pos)
```

```
{
```

```
    struct scull_dev *dev = filp->private_data;
```

```
    if (mutex_lock_interruptible(&dev->lock))  
        return -ERESTARTSYS;
```

```
    if (*f_pos >= dev->size)  
        goto out;
```

```
out:
```

```
    mutex_unlock(&dev->lock);  
    return retval;
```

```
}
```

- Use -ERESTARTSYS when retrying the operation is the right thing to do (undo any user visible changes)
- Use -EINTR when you can't undo previous changes
- Balance each lock with unlock (including error case)

Reader/Writer Semaphores

- How many threads can read a non-atomic variable concurrently when writes to the variable are blocked?
 - Infinite
- How many threads can read a non-atomic variable concurrently when writes to the variable are occurring?
 - 0
- Does every thread need an exclusive lock to access a variable?
 - Only writer threads need exclusive access.
 - An infinite number of reader threads can access as long as writer threads are blocked.

Reader/Writer Semaphores

as follows:

```
void init_rwsem(struct rw_semaphore *sem);
```

```
void down_read(struct rw_semaphore *sem);  
int down_read_trylock(struct rw_semaphore *sem);  
void up_read(struct rw_semaphore *sem);
```

```
void down_write(struct rw_semaphore *sem);  
int down_write_trylock(struct rw_semaphore *sem);  
void up_write(struct rw_semaphore *sem);  
void downgrade_write(struct rw_semaphore *sem);
```

- `down_write_trylock/down_read_trylock` return 1 instead of 0 on success (different from convention)
- `downgrade_write` converts a write lock to a read lock
- Writers get priority
- Use when write access is rarely required, held briefly.

Completions

- Wait for some activity to complete
- Can use a semaphore for this.
 - Semaphores are optimized for “available” case
- Completions designed for the “not available” case
 - Will be the case when you just started a thread and want to wait for it to complete.
- `complete/complete_all()` support use in interrupt handlers.

Completions

- “Any time you think of using `yield()` or some quirky `msleep(1)` loop to allow something else to proceed, you probably want to look into using one of the `wait_for_completion*()` calls and `complete()` instead.”

```
init_completion(&dynamic_object->done);
```

```
static DECLARE_COMPLETION(setup_done);  
DECLARE_COMPLETION(setup_done);
```

```
DECLARE_COMPLETION_ONSTACK(setup_done)
```

Define and initialize the completion

Pass it to a worker that runs in a different thread

Wait for the worker to signal completion

Worker Thread

CPU#1

```
struct completion setup_done;  
  
init_completion(&setup_done);  
initialize_work(...,&setup_done,...);  
  
/* run non-dependent code */  
  
wait_for_completion(&setup_done);
```

CPU#2

```
/* do setup */  
  
complete(setup_done);
```

Completions

- Wait for completion variants

```
unsigned long wait_for_completion_timeout(struct completion *done, unsigned long timeout)
```

```
int wait_for_completion_interruptible(struct completion *done)
```

```
long wait_for_completion_interruptible_timeout(struct completion *done, unsigned long timeout)
```

- Completion Variants

- complete_all - current *and* future waiters

```
void complete(struct completion *done)
```

```
void complete_all(struct completion *done)
```