

# Additional Kernel Debug Strategies

**Advanced Embedded Software  
Development**  
with **Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

Understand kernel debug strategies for driver development

# Using an interactive debugger

- Unfortunately not as easy as with user space programs
- Generally time consuming/should be avoided
  - Compile with CONFIG\_DEBUG\_INFO
- Linus is not a fan of interactive debuggers

```
Date: Wed, 6 Sep 2000 12:52:29 -0700 (PDT)
From: Linus Torvalds <torvalds@transmeta.com>
To: Tigran Aivazian <tigran@veritas.com>
Subject: Re: Availability of kdb
```

On Wed, 6 Sep 2000, Tigran Aivazian wrote:

```
>
> very nice monologue, thanks. It would be great to know Linus' opinion. I
> mean, I knew Linus' opinion of some years' ago but perhaps it changed? He
> is a living being and not some set of rules written in stone so perhaps
> current stability/highquality of kdb suggests to Linus that it may be
> (just maybe) acceptable into official tree?
```

```
I don't like debuggers. Never have, probably never will. I use gdb all the
time, but I tend to use it not as a debugger, but as a disassembler on
```

# Using kdb and kgdb

- Both are now built into the kernel with CONFIG\_KGDB\_KDB and associated configuration (book is out of date on this).
- kdb is the shell you can interact with on the system
  - Typically uses the serial port for communication
  - Entered automatically after oops
- gdb on a development machine interacts with kdb/kgdb on the target, typically through a serial port
- Use SysRq->G to interrupt and enter kgdb mode

# Other Debugging Tools

- User Mode Linux
  - No longer maintained/updated
- Linux Trace Toolkit
  - Traces events in the kernel for debugging and performance issues
- Dynamic Probing
  - replaced by SystemTap
  - Script framework to hook into Linux kernel

# debugfs

- A way for developers to make information available to userspace.
  - /proc - fallen out of favor for new development - intended for process information
  - /sysfs - highly organized/restricted content
- “Not intended to serve as a stable ABI to user space”

# Kernel Memory Leaks

## Validation:

1. Your driver should pass the drivertest.sh script provided with the assignment
2. Your qemu instance should pass the sockettest script, this time writing and reading from only the /dev/aesdchar device instead of /var/tmp/aesdsocketdata.
3. Your implementation should not crash or have memory leaks.

- How do you test for memory leaks?
- You can use kmemleak for this
  - make linux-menuconfig from buildroot
    - kernel hacking->kernel debugging
    - kernel hacking->memory debugging->kernel memory leak detector

```

Memory Debugging
bmenus ----> (or empty submenus ---->). Highlighted letters are hotkeys.
, </> for Search. Legend: [*] built-in [ ] excluded <M> module <>

[ ] Extend mmap on extra space for more information on page
[ ] Debug page memory allocations (NEW)
[ ] Poison pages after freeing
[ ] Testcase for the marking rodata read-only
[ ] Debug object operations (NEW)
[ ] SLUB debugging on by default
[ ] Enable SLUB performance statistics
[*] Kernel memory leak detector
(400) Maximum kmemleak early log entries (NEW)
<> Simple test for the kernel memory leak detector (NEW)
[ ] Default kmemleak to off (NEW)
[ ] Stack utilization instrumentation (NEW)
[ ] Debug VM (NEW)
[ ] Debug VM translations (NEW)
[ ] Debug access to per_cpu maps (NEW)
[ ] KASan: runtime memory debugger
  
```

# Kernel Memory Leaks

```
# mount -t debugfs nodev /sys/kernel/debug/
```

To display the details of all the possible scanned memory leaks:

To test a critical section on demand with a clean kmemleak do:

```
# echo clear > /sys/kernel/debug/kmemleak  
... test your kernel or modules ...  
# echo scan > /sys/kernel/debug/kmemleak
```

Then as usual to get your report with:

```
# cat /sys/kernel/debug/kmemleak
```



# bash -x (or sh -x)

- example `bash -x ./drivertest.sh > /tmp/bash_result.txt 2>&` on unimplemented driver

```
+ cat /tmp/fileyqlkKR
+ rc=-1
+ echo write11
./drivertest.sh: line 52: echo: write error: Cannot allocate memory
++ tempfile
+ expected_file_2_to_11=/tmp/filePYDw8P
+ cat
+ cat /dev/aesdchar
+ echo 'The output should show writes 2-11 in order'
The output should show writes 2-11 in order
+ cat /tmp/fileyqlkKR
+ check_output /tmp/fileyqlkKR /tmp/filePYDw8P
+ local read_file=/tmp/fileyqlkKR
+ local expected_file=/tmp/filePYDw8P
+ diff /tmp/fileyqlkKR /tmp/filePYDw8P
```

-x

Print a trace of simple commands, for commands, case commands, select commands, and arithmetic for commands and their arguments or associated word lists after they are expanded and before they are executed. The value of the PS4 variable is expanded and the resultant value is printed before the command and its expanded arguments.

# strace

- example `strace -o /tmp/strace.txt ./drivertest.sh` on unimplemented driver

```
echo "write1" > ${device}
```

```

openat(AT_FDCWD, "/dev/aesdchar", O_WRONLY|O_CREAT|O_TRUNC, 0666) = 3
fcntl(1, F_GETFD) = 0
fcntl(1, F_DUPFD, 10) = 10
fcntl(1, F_GETFD) = 0
fcntl(10, F_SETFD, FD_CLOEXEC) = 0
dup2(3, 1) = 1
close(3) = 0
fstat(1, {st_mode=S_IFCHR|0664, st_rdev=makedev(240, 0), ...}) = 0
ioctl(1, TCGETS, 0x7fff27f66430) = -1 ENOTTY (Inappropriate ioctl for device)
write(1, "write1\n", 7) = -1 ENOMEM (Cannot allocate memory)

```

Open device as fd 3

Direct stdout to fd 3 (our driver device endpoint)

Write to stdout (directed to our device driver)  
Why did we get ENOMEM?

# strace

- example `strace -o /tmp/strace-aesdsocket.txt -f ./aesdsocket` on working driver

```
16943 openat(AT_FDCWD, "/dev/aesdchar", O_WRONLY|O_CREAT|O_APPEND, 0666) = 7
16943 lseek(7, 0, SEEK_END)          = -1 ESPIPE (Illegal seek)
16943 write(1, "Received String: abcdefg\n", 25) = 25
16943 fstat(7, {st_mode=S_IFCHR|0664, st_rdev=makedev(240, 0), ...}) = 0
16943 ioctl(7, TCGETS, 0x7f3bcb8f1c50) = -1 ENOTTY (Inappropriate ioctl for device)
16943 write(7, "abcdefg\n", 8)          = 8
16943 close(7)                        = 0
```

**-f**

Trace child processes as they are created by currently traced processes as a result of the **fork(2)** system call.