

ioctl

Advanced Embedded Software Development with **Dan Walkes**



University of Colorado **Boulder**

Learning objectives:

Understand a kernel ioctl and its purpose.

Understand how to write an ioctl handler for a kernel driver.

ioctl

- ioctl is a system call on a file descriptor which is passed to the driver.
- Identifies a command to be performed and another argument (typically a pointer)
- Requires another program (or python script?) to issue through system call.

ioctl

- Common interface used for device control
- Request for a device other than read/write, for example
 - Lock a door
 - eject media
 - report information
- User space prototype:

```
#include <sys/ioctl.h>

int ioctl(int fd, unsigned long request, ...);
```

ioctl

```
#include <sys/ioctl.h>

int ioctl(int fd, unsigned long request, ...);
```

- What do the dots mean?
 - Typically represent varargs
 - In this case, are a single optional argument `argp`
 - `request` argument describes `argp` use

The second argument is a device-dependent request code. The third argument is an untyped pointer to memory. It's traditionally `char *argp` (from the days before `void *` was valid C), and will be so named for this discussion.


ioctl

- Easiest and most straightforward choice for device operations
- Has fallen “out of favor among kernel developers”
 - Unstructured system calls
 - difficult to audit
 - often not well documented
- Alternatives
 - embedding commands into the data stream
 - Using virtual filesystems like sysfs

ioctl

- file operations table signature
 - Has changed in 2005 from the type mentioned in the book
 - inode argument referenced in the book available at filp->f_inode

```
int (*ioctl) (struct inode *inode, struct file *filp,
              unsigned int cmd, unsigned long arg);
```



```
long (*unlocked_ioctl) (struct file *filp, unsigned int cmd,
                        unsigned long arg);
```

```
struct file {
```

```
    struct inode *f_inode;
```

```
}
```

```
struct file_operations {
```

```
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *,
                    ssize_t (*write) (struct file *, const char __u
```

```
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
```

```
}
```

Linux Device Drivers 3rd Edition Chapter 6

<https://elixir.bootlin.com/linux/v5.3.1/source/include/linux/fs.h#L1789>

<https://lwn.net/Articles/119652/>

<https://elixir.bootlin.com/linux/v5.3.1/source/include/linux/fs.h#L928>

ioctl

- “The optional arg argument is passed in the form of an unsigned long, regardless of whether it was given by the user as an integer or pointer”
- What is the problem with this approach?
 - Assumes `sizeof(unsigned long) = sizeof(void *)`
- New implementation includes `compat_ioctl` to handle differences in arg size between 64 and 32 bit
 - Avoids breaking 32 bit programs running on 64 bit kernels with assumptions about arg size

```
long (*compat_ioctl) (struct file *filp, unsigned int cmd,  
                     unsigned long arg);
```


ioctl example

- Switch statement based on cmd argument

```
long scull_ioctl(struct file *filp, unsigned int cmd, unsigned long arg)
```

```
{
```

```
switch(cmd) {
```

```
case SCULL_IOCTLRESET:
```

```
    scull_quantum = SCULL_QUANTUM;
```

```
    scull_qset = SCULL_QSET;
```

```
    break;
```

```
case SCULL_IOCSEQUANTUM: /* Set: arg points to the value */
```

```
    if (!capable(CAP_SYS_ADMIN))
```

```
        return -EPERM;
```

```
    retval = __get_user(scull_quantum, (int __user *)arg);
```

```
    break;
```

```
#define SCULL_IOCTLRESET    _IO(SCULL_IOC_MAGIC, 0)
```

```
#define SCULL_IOCSEQUANTUM _IOW(SCULL_IOC_MAGIC, 1, int)
```

Selecting ioctl cmd numbers

- Macros help us generate unique cmd codes for each device driver using “magic” numbers specific to driver
 - Avoids issuing the correct command to the wrong device

```
/* Use 'k' as magic number */
```

```
#define SCULL_IOC_MAGIC 'k'
```

```
#define SCULL_IOCRESET _IO(SCULL_IOC_MAGIC, 0)
```

Code	Seq# (hex)	Include File	Comments
'k'	00-0F	linux/spi/spidev.h	conflict!
'k'	00-05	video/kyro.h	conflict!
'k'	10-17	linux/hsi/hsi_char.h	HSI character device

The first argument to _IO, _IOW, _IOR, or _IOWR is an identifying letter or number from the table below. Because of the large number of drivers, many drivers share a partial letter with other drivers.

Using _IO macros

- First argument - magic number
- Second arg - sequence (command) number
- 3rd arg - type of data transferred

```
#define _IOW(type,nr,size)      _IOC(_IOC_WRITE,(type),(nr),(_IOC_TYPECHECK(size)))
```

```
#define _IOC(dir,type,nr,size) \
    (((dir) << _IOC_DIRSHIFT) | \
     ((type) << _IOC_TYPESHIFT) | \
     ((nr) << _IOC_NRSHIFT) | \
     ((size) << _IOC_SIZESHIFT))
```

```
#define SCULL_IOCRESET      _IO(SCULL_IOC_MAGIC, 0)
```

```
#define SCULL_IOCSEQUANTUM _IOW(SCULL_IOC_MAGIC, 1, int)
```

_IO	an	ioctl with no parameters
_IOW	an	ioctl with write parameters (copy_from_user)
_IOR	an	ioctl with read parameters (copy_to_user)
_IOWR	an	ioctl with both write and read parameters.

Using _IO macros

- _IOC_READ (_IOR) - transfer from kernel to user space
- _IOC_WRITE (_IOW)- transfer from user to kernel space

```
#define _IOW(type,nr,size) _IOC(_IOC_WRITE,(type),(nr),(_IOC_TYPECHECK(size)))
```

```
#define _IOC(dir,type,nr,size) \
    (((dir) << _IOC_DIRSHIFT) | \
     ((type) << _IOC_TYPESHIFT) | \
     ((nr) << _IOC_NRSHIFT) | \
     ((size) << _IOC_SIZESHIFT))
```

```
#define SCULL_IOCRESET _IO(SCULL_IOC_MAGIC, 0)
```

```
#define SCULL_IOCSEQUANTUM _IOW(SCULL_IOC_MAGIC, 1, int)
```

_IO	an	ioctl with no parameters
_IOW	an	ioctl with write parameters (copy_from_user)
_IOR	an	ioctl with read parameters (copy_to_user)
_IOWR	an	ioctl with both write and read parameters.

Predefined commands

- Some ioctls are recognized by the kernel, decoded before passing to your driver
- Avoid magic type “T” to avoid these

Code	Seq# (hex)	Include File	Comments
'T'	all	arch/x86/include/asm/ioctls.h	conflict!

write() as ioctl alternative

- Write control sequences to the device
 - Doesn't require a program (ioctl) to implement
 - echo "start" > /dev/yourdevice
 - Especially useful for devices which don't transfer data
 - May make the driver more complex in some cases
 - if so ioctl or sysfs may be a better solution