

Kernel Sleep Continued

**Advanced Embedded Software
Development**
with **Dan Walkes**



University of Colorado **Boulder**

Learning objectives:

Understand scull pipe wakeup kernel example from driver sources.

Understand nonblocking I/O implementation as an alternative to sleep.

Scull Pipe Sleep/Wakeup Example

```
static ssize_t scull_p_read (struct file *filp, char __user *buf, size_t count,
                           loff_t *f_pos)
{
    while (dev->rp == dev->wp) { /* nothing to read */
        mutex_unlock(&dev->lock); /* release the lock */
        if (filp->f_flags & O_NONBLOCK)
            return -EAGAIN;
        PDEBUG("\%s\% reading: going to sleep\n", current->comm);
        if (wait_event_interruptible(dev->inq, (dev->rp != dev->wp)))
            return -ERESTARTSYS; /* signal: tell the fs layer to handle it */
        /* otherwise loop, but first reacquire the lock */
        if (mutex_lock_interruptible(&dev->lock))
            return -ERESTARTSYS;
    }

    static ssize_t scull_p_write(struct file *filp, const char __user *buf, size_t count,
                                loff_t *f_pos)
    {
        -- -- --

        dev->wp += count;
        if (dev->wp == dev->end)
            dev->wp = dev->buffer; /* wrapped */
        mutex_unlock(&dev->lock);

        /* finally, awake any reader */
        wake_up_interruptible(&dev->inq); /* blocked in read() and select() */
    }
}
```



Why release the lock here?

- Don't sleep with lock held

```
ecen5013@ecen5013-VirtualBox:~/l3dd3$ cat /dev/scullpipe
```

```
echo "hello world!" >/dev/scullpipe
```

```
ecen5013@ecen5013-VirtualBox:~/l3dd3$ cat /dev/scullpipe
hello world!
```

Blocking I/O Buffering

- Output and Input buffers are often useful for handling blocking I/O on real devices.
 - This means any blocking is on access to the buffer rather than access to the device
 - Especially large performance benefit when device interaction is much slower than memory access

How a process sleeps

- Task states:
 - TASK_RUNNING - able to run, may not be executing
 - TASK_INTERRUPTIBLE and TASK_UNINTERRUPTIBLE - two types of sleep
- schedule()
 - function which runs the scheduler, allows other process to be scheduled (effectively puts the process to sleep)

How a process sleeps

```
wait_event_interruptible(wq, flag != 0);
```

```
#define __wait_event_interruptible(wq_head, condition)
__wait_event(wq_head, condition, TASK_INTERRUPTIBLE, 0, 0,
schedule())
```

```
long prepare_to_wait_event(struct wait_queue_head *wq_head, struct wait_queue_entry *wq_entry, int state)
{
```

```
spin_lock_irqsave(&wq_head->lock, flags);
```

```
set_current_state(state);
```

```
spin_unlock_irqrestore(&wq_head->lock, flags);
```

```
}
```

```
void finish_wait(struct wait_queue_head *wq_head, struct wait_queue_entry *wq_entry)
```

```
{
```

```
__set_current_state(TASK_RUNNING);
```

```
}
```

```
#define __wait_event(wq_head, condition, state, exclusive, ret, cmd)
({
    __label__ __out;
    struct wait_queue_entry __wq_entry;
    long __ret = ret; /* explicit shadow */

    init_wait_entry(&__wq_entry, exclusive ? WQ_FLAG_EXCLUSIVE : 0);
    ; {
        long __int = prepare_to_wait_event(&wq_head, &__wq_entry, state);

        if (condition)
            break;

        if (__wait_is_interruptible(state) && __int) {
            __ret = __int;
            goto __out;
        }

        cmd;

    }
    finish_wait(&wq_head, &__wq_entry);
    __out: __ret;
})
```



avoid waiting too long by checking the condition before sleeping

Linux Device Drivers 3rd Edition Chapter 6

<https://elixir.bootlin.com/linux/v5.3.1/source/include/linux/sched.h#L74>

<https://elixir.bootlin.com/linux/v5.3.1/source/include/linux/wait.h#L455>

<https://elixir.bootlin.com/linux/v5.3.1/source/kernel/sched/wait.c#L258>

Nonblocking I/O

- What if process requests O_NONBLOCK on open()?
 - O_NONBLOCK will be set in filp->f_flags

```
static int scull_w_open(struct inode *inode, struct file *filp)
{
    struct scull_dev *dev = &scull_w_device; /* device information */

    spin_lock(&scull_w_lock);
    while (! scull_w_available()) {
        spin_unlock(&scull_w_lock);
        if (filp->f_flags & O_NONBLOCK) return -EAGAIN;
        if (wait_event_interruptible (scull_w_wait, scull_w_available()))
            return -ERESTARTSYS; /* tell the fs layer to handle it */
        spin_lock(&scull_w_lock);
    }
}
```

Testing Non-blocking IO

- Book references a nbtest example which demonstrates non-blocking operation

```
int main(int argc, char **argv)
{
    int delay = 1, n, m = 0;

    if (argc > 1)
        delay=atoi(argv[1]);
    fcntl(0, F_SETFL, fcntl(0,F_GETFL) | O_NONBLOCK); /* stdin */
    fcntl(1, F_SETFL, fcntl(1,F_GETFL) | O_NONBLOCK); /* stdout */

    while (1) {
        n = read(0, buffer, 4096);
        if (n >= 0)
            m = write(1, buffer, n);
        if ((n < 0 || m < 0) && (errno != EAGAIN))
            break;
        sleep(delay);
    }
    perror(n < 0 ? "stdin" : "stdout");
    exit(1);
}
```


Testing Nonblocking IO

- Use misc-progs and nonblock.sh test script after loading the driver you want to test

```
./misc-progs/test/nonblock.sh -i /dev/scullpipe -s
```

```
Reading content of infile through non blocking test to /dev/scullpipe
strace ../nbtest 1 > /dev/scullpipe < infile
```

```
fcntl(0, F_GETFL)          = 0x8000 (flags O_RDONLY|O_LARGEFILE)
fcntl(0, F_SETFL, O_RDONLY|O_NONBLOCK|O_LARGEFILE) = 0
fcntl(1, F_GETFL)          = 0x8001 (flags O_WRONLY|O_LARGEFILE)
fcntl(1, F_SETFL, O_WRONLY|O_NONBLOCK|O_LARGEFILE) = 0
read(0, "", 4096)          = 0
write(1, "", 0)             = 0
nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffe04147b50) = 0
read(0, "", 4096)          = 0
write(1, "", 0)             = 0
```

```
echo "Hello World AESD!" >> misc-progs/test/infile
```

```
ecen5013@ecen5013-VirtualBox:~/l3dd3$ tail -f misc-progs/test/outfile
```

```
Hello World AESD!
```

```
./misc-progs/test/nonblock.sh -o /dev/scullpipe -s
```

```
Reading content of /dev/scullpipe through non blocking test to outfile
strace ../nbtest 1 > outfile < /dev/scullpipe
```

```
fcntl(0, F_GETFL)          = 0x8000 (flags O_RDONLY|O_LARGEFILE)
fcntl(0, F_SETFL, O_RDONLY|O_NONBLOCK|O_LARGEFILE) = 0
fcntl(1, F_GETFL)          = 0x8001 (flags O_WRONLY|O_LARGEFILE)
fcntl(1, F_SETFL, O_WRONLY|O_NONBLOCK|O_LARGEFILE) = 0
read(0, 0x5608ff3d6040, 4096) = -1 EAGAIN (Resource temporarily unavailable)
nanosleep({tv_sec=1, tv_nsec=0}, 0x7fff0e4b82f0) = 0
read(0, 0x5608ff3d6040, 4096) = -1 EAGAIN (Resource temporarily unavailable)
nanosleep({tv_sec=1, tv_nsec=0}, 0x7fff0e4b82f0) = 0
read(0, 0x5608ff3d6040, 4096) = -1 EAGAIN (Resource temporarily unavailable)
```

Testing Non-blocking IO

- Writing to misc-progs/test/infile results in output on misc-progs/test/outfile through non-blocking pipe

```

echo "Hello World AESD!" >> misc-progs/test/infile

nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffe1b6e52c0) = 0
read(0, "", 4096) = 0
write(1, "", 0) = 0
nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffe1b6e52c0) = 0
read(0, "Hello World AESD!\n", 4096) = 18
write(1, "Hello World AESD!\n", 18) = 18
nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffe1b6e52c0) = 0
read(0, "", 4096) = 0
write(1, "", 0) = 0

ecen5013@ecen5013-VirtualBox:~/ltd3$ tail -f misc-progs/test/outfile
Hello World AESD!

nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffe1b6e52c0) = 0
read(0, 0x55d07d11c040, 4096) = -1 EAGAIN (Resource temporarily unavailable)
nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffc7dcd9310) = 0
read(0, "Hello World AESD!\n", 4096) = 18
write(1, "Hello World AESD!\n", 18) = 18
nanosleep({tv_sec=1, tv_nsec=0}, 0x7ffc7dcd9310) = 0
read(0, 0x55d07d11c040, 4096) = -1 EAGAIN (Resource temporarily unavailable)

Sep 28 23:01:56 ecen5013-VirtualBox kernel: [36073.766929] scull: "nbtest" did write 0 bytes
Sep 28 23:01:57 ecen5013-VirtualBox kernel: [36074.771680] scull: Going to accept 18 bytes to 0000000044f052ea from 00000000472b6a74
Sep 28 23:01:57 ecen5013-VirtualBox kernel: [36074.771683] scull: "nbtest" did write 18 bytes
Sep 28 23:01:58 ecen5013-VirtualBox kernel: [36075.359357] scull: "nbtest" did read 18 bytes

```