

# Poll and Select

## **Advanced Embedded Software Development**

with **Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

Understand use of poll and select in kernel drivers.

Understand seeking and associated read/write/open function implementation.

# poll and select

- poll and select determine whether a device/file can be read without blocking or wait for a file descriptor to become ready
- Support implemented by poll method in device driver

```
struct file_operations {
```

```
    poll_t (*poll) (struct file *, struct poll_table_struct *);
```

Linux Device Drivers 3rd Edition Chapter 6

<https://elixir.bootlin.com/linux/v5.3.1/source/include/linux/fs.h#L1789>

<https://github.com/cu-ecen-5013/ldd3/blob/latest-lecture-source/scull/pipe.c>

<http://man7.org/linux/man-pages/man2/poll.2.html>

<http://man7.org/linux/man-pages/man2/select.2.html>

```
/*
 * The file operations for the pipe device
 * (some are overlaid with bare scull)
 */
struct file_operations scull_pipe_fops = {
    .owner =      THIS_MODULE,
    .llseek =     no_llseek,
    .read =       scull_p_read,
    .write =      scull_p_write,
    .poll =       scull_p_poll,
    .unlocked_ioctl = scull_ioctl,
    .open =       scull_p_open,
    .release =    scull_p_release,
    .fsync =      scull_p_fsync,
};
```

# poll and select

```
static unsigned int scull_p_poll(struct file *filp, poll_table *wait)
{
    struct scull_pipe *dev = filp->private_data;
    unsigned int mask = 0;

    /*
     * The buffer is circular; it is considered full
     * if "wp" is right behind "rp" and empty if the
     * two are equal.
     */
    mutex_lock(&dev->lock);
    poll_wait(filp, &dev->inq, wait);
    poll_wait(filp, &dev->outq, wait);
    if (dev->rp != dev->wp)
        mask |= POLLIN | POLLRDNORM; /* readable */
    if (spacefree(dev))
        mask |= POLLOUT | POLLWRNORM; /* writable */
    mutex_unlock(&dev->lock);
    return mask;
}
```

```
struct scull_pipe {
    wait_queue_head_t inq, outq; /* read and write queues */
};
```

- Call poll\_wait on any wait queues which could indicate poll status changes
  - Kernel will wait on these as necessary
- Return bit mask describing currently available operations (read or write)

# Read Rules

- When data is available in the input buffer:
  - read should return immediately with at least 1 byte
  - poll should return `POLLIN|POLLRDNORM`
- When no data in the input buffer
  - read
    - should block until one byte is there OR
    - if `O_NONBLOCK` return with value of `-EAGAIN`
  - poll should report unreadable (read flags all 0)

# Read Rules

- At end of file
  - read should return immediately with value of 0
  - poll should report POLLHUP

# Write Rules

- When space is available in the output buffer:
  - write should return without delay, accepting at least one byte
  - poll should report POLLOUT | POLLWRNORM
- When the output buffer is full
  - write
    - should block until space is freed or
    - if O\_NONBLOCK is set return -EAGAIN
  - poll should report file is not writable (write flags all 0)

# Write Rules

- Never make a write call wait for data transmission (transfer from output buffer to device), **even if `O_NONBLOCK` is not set.**
  - To ensure transmissions complete to the device, the driver must provide `fsync`



# Seeking on a Device

- Default lseek just sets filp->f\_pos
- llseek file operation can be implemented if the seek a custom operation for the device.
- Call nonseekable\_open in your open function if seek doesn't make sense for your device
  - data flow rather than data area
    - serial port
    - keyboard

# Access Control from open()

- Single-Open
  - Only one process can open at a time
  - scullsingle example - Obtain atomic in open, release in release()
- Single User
  - Compare process uid in open(), return -EBUSY when in use
- Alternative to returning -EBUSY - block in open()