

# Assignment 9

## Overview

**Advanced Embedded Linux  
Development**  
with **Dan Walkes**



University of Colorado **Boulder**

## **Learning objectives:**

Understand Assignment 9

implementation/ioctl implementation

# f\_pos member



```
loff_t f_pos;

#if defined(__GNUC__)
typedef __kernel_loff_t loff_t;
#endif

typedef long long __kernel_loff_t;
```

- The file structure f\_pos member:
  - References “The current reading or writing position”
    - 64 bit signed offset representing a byte offset in the file.

# f\_pos member

- Passed as an input/output argument to read/write
  - Always 0 in write when using echo “blah” > /dev/aesdchar
    - Why?
    - > is an overwrite of the file, starts at offset 0

```
ssize_t aesd_write(struct file *filp, const char __user *buf, size_t count,  
                  loff_t *f_pos)  
{
```

```
ssize_t aesd_read(struct file *filp, char __user *buf, size_t count,  
                 loff_t *f_pos)
```

# f\_pos member

- Also included in struct file

```
struct file {
```

```
    loff_t
```

```
    f_pos;
```

```
ssize_t aead_write(struct file *filp, const char __user *buf, size_t count,  
                  loff_t *f_pos)
```

```
ssize_t aead_read(struct file *filp, char __user *buf, size_t count,  
                 loff_t *f_pos)
```

- Why include in two places?
  - passed \*f\_pos is modified by read/write implementation
  - filp->f\_pos is modified by underlying shared drivers based on read/write

# Seek Implementation

```
loff_t f_pos;
```

The current reading or writing position. `loff_t` is a 64-bit value on all platforms (long long in `gcc` terminology). The driver can read this value if it needs to know the current position in the file but should not normally change it; *read* and *write* should update a position using the pointer they receive as the last argument instead of acting on `filp->f_pos` directly. The one exception to this rule is in the *llseek* method, the purpose of which is to change the file position.

- We will implement `llseek` in assignment 9
- We will also modify `filp->f_pos` in an `ioctl` (wouldn't typically be done this way)

# Reading Positionally

- **SEEK\_SET**
  - Use the specified offset as the file position.
- **SEEK\_CUR**
  - Increment or decrement file position
- **SEEK\_END**
  - Use EOF as file position

NAME  
lseek - reposition read/write file offset

SYNOPSIS  
#include <sys/types.h>  
#include <unistd.h>  
  
off\_t lseek(int fd, off\_t offset, int whence);

DESCRIPTION  
lseek() repositions the file offset of the open file description associated with the file descriptor fd to the argument offset according to the directive whence as follows:

SEEK\_SET  
The file offset is set to offset bytes.

SEEK\_CUR  
The file offset is set to its current location plus offset bytes.

SEEK\_END  
The file offset is set to the size of the file plus offset bytes.

# Seek Implementation

## The llseek Implementation

The `llseek` method implements the `lseek` and `llseek` system calls. We have already stated that if the `llseek` method is missing from the device's operations, the default implementation in the kernel performs seeks by modifying `filp->f_pos`, the current reading/writing position within the file. Please note that for the `lseek` system call to work correctly, the `read` and `write` methods must cooperate by using and updating the offset item they receive as an argument.

- Why two different system calls, `lseek` and `_llseek`?
  - `llseek` is guaranteed to support long long offset sizes

```
off_t lseek(int fd, off_t offset, int whence);
```

```
int _llseek(unsigned int fd, unsigned long offset_high,  
            unsigned long offset_low, loff_t *result,  
            unsigned int whence);
```



# lseek driver implementation

```
struct file_operations {
```

```
    loff_t (*llseek) (struct file *, loff_t, int);  
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);  
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
```

# lseek driver implementation

- “If the lseek method is missing from the device’s operations, the default implementation in the kernel performs seek by modifying `filp->fpos`.”
- “For the lseek system call to work correctly the read and write methods must cooperate by using and updating the offset item they receive as an argument”

# lseek driver implementation

- “If the lseek method is missing from the device’s operations, the default implementation in the kernel performs seek by modifying filp->fpos.”

```
extern loff_t fixed_size_llseek(struct file *file, loff_t offset,
                               int whence, loff_t size);
```

- Several wrappers around generic llseek which seek for you which you can call from your llseek method
- This one uses a size you provide
- What should we use for size?
  - The total size of all content of the circular buffer

```
/**
 * fixed_size_llseek - llseek implementation for fixed-sized devices
 * @file:      file structure to seek on
 * @offset:    file offset to seek to
 * @whence:    type of seek
 * @size:      size of the file
 */
loff_t fixed_size_llseek(struct file *file, loff_t offset, int whence, loff_t size)
{
    switch (whence) {
        case SEEK_SET: case SEEK_CUR: case SEEK_END:
            return generic_file_llseek_size(file, offset, whence,
                                             size, size);

        default:
            return -EINVAL;
    }
}
EXPORT_SYMBOL(fixed_size_llseek);
```

# Ilseek Assignment 9 Options

1. Leave Ilseek null and use the default Ilseek
  1. Would require supporting seek in write() (not assuming every write appends) which isn't an assignment requirement - I have not experimented with this myself
2. Add your own Ilseek function, with locking **and logging**, but use `fixed_size_Ilseek` for logic.
3. Implement your own Ilseek function separate from `fixed_size_Ilseek` handling each of the “whence” cases

I suggest option 2

# read/write method and f\_pos

- “For the lseek system call to work correctly the read and write methods must cooperate by using and updating the offset item they receive as an argument”
  - read function:
    - Must set `*f_pos` to `*f_pos + retcount` where `retcount` is the number of bytes read
  - write function:
    - Must set `*f_pos` to `*f_pos + retcount` where `retcount` is the number of bytes written

# ioctl implementation

- Add an `aesd_ioctl.h` file you can share with your `aesdsocket` implementation
- See provided `aesd_ioctl.h` file at [https://github.com/cu-ecen-aeld/aesd-assignments/blob/assignment9/aesd-char-driver/aesd\\_ioctl.h](https://github.com/cu-ecen-aeld/aesd-assignments/blob/assignment9/aesd-char-driver/aesd_ioctl.h)

# ioctl user space implementation -

## aesdsocket

```
struct aed_seekto seekto;
seekto.write_cmd = write_cmd;
seekto.write_cmd_offset = offset;
int result_ret = ioctl(fd,AESDCHAR_IOCSEEKTO,&seekto);
```

```
struct aed_seekto {
    /**
     * The zero referenced write command to seek into
     */
    uint32_t write_cmd;
    /**
     * The zero referenced offset within the write
     */
    uint32_t write_cmd_offset;
};
```

- #include the aed\_ioctl.h file
- Use the ioctl command with fd representing the driver
- Pass the filled in structure to the driver via ioctl

```
#include <sys/ioctl.h>

int ioctl(int fd, unsigned long request, ...);
```

# ioctl implementation - aesdaocket

- My aesdaocket uses a FILE\* to access my driver. How do I get the fd used for ioctl?
  - Use fileno()

```
#include <stdio.h>

int fileno(FILE *stream);
```

```
#include <sys/ioctl.h>

int ioctl(int fd, unsigned long request, ...);
```



# ioctl implementation - driver

```
case AESDCHAR_IOCSEEKTO:
{
    struct aese_seekto seekto;
    if( copy_from_user(&seekto, (const void __user *)arg, sizeof(seekto)) != 0 ) {
        retval = EFAULT;
    } else {
        retval = aese_adjust_file_offset(filp,seekto.write_cmd,seekto.write_cmd_offset);
    }
    break;
}
```

- #include the aese\_ioctl.h file
- Use copy\_from\_user to obtain the value from userspace

# Adjusting the file offset from ioctl



```
/**
 * Adjust the file offset (f_pos) parameter of @param filp based on the location specified by
 * @param write_cmd (the zero referenced command to locate)
 * and @param write_cmd_offset (the zero referenced offset into the command)
 * @return 0 if successful, negative if error occurred:
 *     -ERESTARTSYS if mutex could not be obtained
 *     -EINVAL if write command or write_cmd_offset was out of range
 */
static long aesd_adjust_file_offset(struct file *filp, unsigned int write_cmd, unsigned int write_cmd_offset)
```

- Check for valid write\_cmd and write\_cmd\_offset values
- Calculate the start offset to write\_cmd
- Add write\_cmd\_offset
- Save as filp->f\_pos

# Adjusting the file offset from ioctl



```
/**
 * Adjust the file offset (f_pos) parameter of @param filp based on the location specified by
 * @param write_cmd (the zero referenced command to locate)
 * and @param write_cmd_offset (the zero referenced offset into the command)
 * @return 0 if successful, negative if error occurred:
 *     -ERESTARTSYS if mutex could not be obtained
 *     -EINVAL if write command or write_cmd_offset was out of range
 */
static long aesd_adjust_file_offset(struct file *filp, unsigned int write_cmd, unsigned int write_cmd_offset)
```

- Check for valid write\_cmd and write\_cmd\_offset values
- When would values be invalid?
  - haven't written this command yet
  - out of range cmd (11)
  - write\_cmd\_offset is  $\geq$  size of command

# Adjusting the file offset from ioctl



```
/**
 * Adjust the file offset (f_pos) parameter of @param filp based on the location specified by
 * @param write_cmd (the zero referenced command to locate)
 * and @param write_cmd_offset (the zero referenced offset into the command)
 * @return 0 if successful, negative if error occurred:
 *     -ERESTARTSYS if mutex could not be obtained
 *     -EINVAL if write command or write_cmd_offset was out of range
 */
static long aesd_adjust_file_offset(struct file *filp, unsigned int write_cmd, unsigned int write_cmd_offset)
```

- Calculate the start offset to write\_cmd
  - Add length of each write between the output pointer and write\_cmd
- Add the write\_cmd\_offset

# New Test Scripts

- drivertest.sh and sockettest.sh scripts should continue to pass with the changes in this assignment.
- New scripts [assignment9/drivertest.sh](#) and [assignment9/sockettest.sh](#) should also succeed.
  - You will need to stop/restart your driver when running the sockettest scripts (since we aren't supporting command delete).

# Seeking on the command line

- Use the dd utility to perform seeking on the command line

```
read_with_seek()
{
    local seek=$1
    local device=$2
    local read_file=$3
    dd if=${device} skip=${seek} of=${read_file} bs=1 > /dev/null 2>&1
}
```

```
device=/dev/aesdchar
```

```
read_file=$(mktemp)
expected_file=$(mktemp)
```

```
read_with_seek 2 ${device} ${read_file}
```

NAME

dd - convert and copy a file

if=FILE

read from FILE instead of stdin

skip=N skip N ibs-sized blocks at start of input

of=FILE

write to FILE instead of stdout

bs=BYTES

read and write up to BYTES bytes at a time (default: 512);

<http://man7.org/linux/man-pages/man1/dd.1.html>

<https://github.com/cu-ecen-5013/aesd-assignments/blob/assignment9/aesd-char-driver/drivertest-assignment-9.sh>